

# Lab 3: TCP Congestion and UDP Socket Programming

50.012 Networks (2018)

Hand out: 26<sup>th</sup> Sept (C1) / 28<sup>th</sup> Sept (C2)

Hand in: 2<sup>nd</sup> Oct (C1) / 4<sup>th</sup> Oct (C2)

## 1 Objectives

- Experiment with buffer sizes, network scheduling, and TCP congestion window
- Develop a simple client-server application based on UDP sockets
- Goal: Understand why larger buffers can be disadvantageous
- Part of this exercise is based on the Stanford CS244 class lab1

## 2 Setup

### 2.1 Lab3 setup

- Download the lab3.zip from eDimension and unpack to some local folder, e.g. ~/lab3/
- After decompressing lab3.zip content, make sure all .sh binaries in lab3 folder are executable. To make binaries executable use the following command:  
`chmod +x *.sh`
- Change into the lab3 folder, and open up another terminal tab with CTRL+SHIFT+T

### 2.2 mininet

- This exercise assumes that you have a running mininet installation
  - It should be installed on the lab machines already
  - On (x)ubuntu, you can install it by running our `install.sh`. Mininet can run on many \*nixes, but NOT on Windows or Mac

#### 1. What is Mininet?

- Mininet is a network emulator (mostly written in Python), that allows you to define network topologies and to connect emulated hosts.
- The emulated hosts run on your OS, and can execute your client/server applications easily.

- You can also easily change link parameters such a packet loss rate, delay, and bandwidth.
- Mininet is mainly intended to experiment with

*Software Defined Networks and OpenFlow*, but is also convenient for this lab session. We might look at OpenFlow later in the term.

- More details at <http://mininet.org>, in particular <http://mininet.org/walkthrough/>
- If you like videos: some basic things of this sheet are also demo'ed in <https://www.youtube.com/watch?v=jmlgXaocwiE>

## 2. Installation and first test

- Start up mininet:

```
sudo ./run.sh
```

- Get familiar with mininet commands using mininet's help system:

```
mininet> help [topic]
```

- Use the following mininet command to spawn two terminals, belonging to two nodes in the emulated network.

```
mininet> xterm h1 h2
```


- Using these terminals, you can run applications or perform ping operations on the emulated hosts
  - To close running commands like `ping -c 100`, try CTRL-C
- Find out the IP addresses of h1 and h2
- Note that `mininet` allows to run a command from its command prompt as if the command is run from any virtual host. For example, if we want to ping h2 from h1 we can run:

```
mininet> h1 ping h2
```

- To leave/close Mininet, try CTRL-D on the main command line


## 3 Queue Length and TCP

In this exercise we will study the dynamics of TCP in home networks. Take a look at the figure below which shows a "typical" home network with a "Home Router" connected to an end host. The "Home Router" is connected via Cable or DSL to a Headend router at the Internet access provider's office. We are going to study what happens when we download data from a remote server to the End Host in this home network.



Problem overview

In a real network it's hard to measure the TCP congestion window  $cwnd$  (because it's private to the Server) and the buffer occupancy (because it's private to the router). Luckily, we have mininet to help allow us to get these values easily.



Topology for Mininet

### 3.1 Simple TCP connection (wget)

- Start up mininet

```
sudo ./run.sh
```

- In the running mininet emulation, test how long it takes to download a webpage on h1 from h2.
- wget's visualization is a bit confusing. You can find the total

time by subtracting start time from end time, or by looking at the time at end of last part line.

```
mininet> h2 wget h1
```

- We discussed the TCP congestion window in the last week, what is your expectation how the  $cwnd$  changes during this HTTP session?
- How big is the router queue/buffer in your opinion?

### 3.2 With parallel load (iperf)

- To see how the dynamics of a long flow differs from a short flow (which never leaves slow-start), we are going to repeat Part 2 for a "streaming video flow". Instead of actually watching videos on your machine, we are going to set up a long-lived high speed TCP connection instead, to emulate a long-lived video flow.

- You can generate long flows using the `iperf` command. `iperf` is a tool for performing network throughput measurements. It can test either TCP or UDP throughput. To perform an `iperf` test the user must establish both a server (to discard traffic) and a client (to generate traffic). We have wrapped `iperf` in a script that you can run as follows:

```
mininet> h1 ./iperf.sh
```

- You can see the throughput of TCP flow from h1 to h2 by running:

```
mininet> h2 tail -f ./iperf-recv.txt
```

- You can quit viewing throughput by pressing CTRL-C.
- Think about the congestion window of the TCP stream again. How will it look for a long-established connection?
- You can also use ping to observe RTT while `iperf` is running

```
mininet> h1 ping -c 100 h2
```

#### The Impact on the Short Flow

- To see how our long-lived `iperf` flow affects our web page download, download the webpage again - while `iperf` is running. Observe how long it takes.

```
mininet> h2 wget h1
```

- Why does the web page take so much longer to download?

### 3.3 Measuring the real cwnd and buffer occupancy values.

We provided a script that lets you measure cwnd and buffer occupancy values (in terms of values for cwnd and buffer occupancy). We are going to re-run a couple of the experiments and plot the real values.

#### 1. Restart Mininet

- Stop and restart Mininet and the monitor script, then re-run the above experiment as follows.

```
mininet> exit
sudo ./run.sh
```

#### 2. Monitor TCP CWND and Buffer Occupancy in Mininet

- In your other terminal tab, type the following giving a name for your experiment.

```
./monitor.sh <EXP_NAME>
```

- In your first tab with the running mininet, start `iperf` again

```
mininet> h1 ./iperf.sh
```

- (wait for 70 seconds ...)

```
mininet> h2 wget h1
```

- Wait for the wget to complete, then stop the python monitor script followed by the instructions on the screen. The cwnd values are saved in <EXP\_NAME>\_tcpprobe.txt and the buffer occupancy in <EXP\_NAME>\_sw0-qlen.txt.

### 3. Plot CWND and Queue Occupancy

- Plot the TCP cwnd and queue occupancy from the output file with the following:

```
./plot_figures.sh <EXP_NAME>
```

Adjust command line parameters to generate the figure you want.

- The script will host a webserver on the machine and you can use the url the script provided to access to your figures.
- Note: the figures will sometimes interpolate where it is not appropriate. If you see a slow decrease of cwnd, then that probably is an artefact of such interpolation. Ignore those segments of the figures.
- If you are unable to see the cwnd, ensure you run wget after you started the monitor.sh script.

## 3.4 Smaller Buffer: from 100 packets to 20 packets

- Stop any running Mininet. This time we will make the buffers 20 packets long instead using:

```
sudo ./run-minq.sh
```

- Repeat the earlier simple tests

```
mininet> h2 wget h1
```

```
mininet> h1 ping -c 10 h2
```

- Did the results change?
- Close mininet, restart the experiment and start monitoring again

```
sudo ./run-minq.sh
```

```
sudo ./monitor.sh <EXP_NAME>
```

- Re-run the experiment:

```
mininet> h1 ./iperf.sh
```

```
mininet> h1 ping -c 30 h2
```

```
mininet> h2 wget h1
```

- What do you think the cwnd and queue occupancy will be like in this case?
- Plot the figure for cwnd and queue occupancy, this time using:

```
./plot_figures_minq.sh <EXP_NAME>
```

- Then again, use the url to see your figures. Why does reducing the queue size reduce the download time for wget?

## 4 UDP Experiments

- In this part, you write a simple UDP client/server system
- We are providing `client.py` and `server.py` scripts for your convenience.
  - We expect that `client.py` can be called with `-r` parameter, and we already provide the code for doing it in the attached client script.
- A helpful tutorial on a similar application: <https://pymotw.com/2/socket/udp.html>

### 4.1 Simple UDP client/server application using sockets

Your UDP server/client should follow these requirements:

- The server will receive incoming datagrams on port 5555
- The exact message format is up to you, it does not need to be efficient
- The client should send messages with a leading segment ID (e.g. part of the UDP payload)
- Your client should be configurable to send a fix amount of data per second (e.g. 1 Mbps), including UDP and IP header
- The server should verify if incoming traffic has missing datagrams, and display a warning if that is the case

### 4.2 Simple UDP connection

- start up mininet

```
sudo ./run.sh
```

- In the running mininet emulation, open a xterm window for h1, and for h2

```
mininet> xterm h1 h2
```

- Start your server application on h2, and the client on h1. The client should now be able to send data to the server.
- Set up your client to send approximately 1.5 Mbps (Megabit per second) of traffic to h2. Do you see dropped packets at the server?
- Set up your client to send a bit more, e.g. 1.6 Mbps. Again, do you see dropped packets?

## 5 What to Hand in

### 5.1 eDimension submission:

- Please provide a writeup (1-2 pages, PDF format with your name) that includes the following information:
  - What is the normal time required to download the webpage on h1 from h2?
  - What was your initial expectation for the congestion window size over time?
  - After starting iperf on h1, did you observe something interesting in the ping RTT?
  - After starting iperf on h1, why does the web page take so much longer to download?
  - Please provide the figures for the first experiment (with qlen 100 and only one queue)
    - \* Please comment on what you can see in the figures
  - Please provide the figures for the second experiment (with qlen 20 and only one queue)
    - \* Please comment on what you can see in the figures, and what is different (and why)
- Submit your server and client python code for the UDP part
  - The client should send approximately 1.5 Mbps
  - No packets should be dropped by default
  - If we increase the packet rate by 10%, we should see dropped packets
  - The server should report if incoming traffic has dropped datagrams (and, optionally, if a new sequence of data is started)
  - You are allowed to collaborate with another student, please specify both names and IDs in the client and server script files.

**NOTE:** Please note that the due dates for submission are 2<sup>nd</sup> Oct for Cohort 1 and 4<sup>th</sup> Oct for Cohort 2. We will look at the timestamps of your submission to evaluate timely / late submission.