

Lec3 – The Application Layer

50.012 Networks

Jit Biswas

Cohort 1: TT7&8 (1.409-10)

Cohort 2: TT24&25 (2.503-4)

Lec3 – The Application Layer

50.012 Networks

Jit Biswas

Cohort 1: TT7&8 (1.409-10)

Cohort 2: TT24&25 (2.503-4)

Introduction

- Today's lecture: The application layer
 - Application layer addressing/ URLs
 - Sockets
 - FTP as an example application layer protocol
 - HTTP as an important application layer protocol
 - P2P Applications

The Application Layer

Overview of the application layer

- Contains protocols that are specific to applications
- These protocols can rely on lower layers to provide routing, error correction, etc.
- Examples application layer protocols:
 - FTP and SMTP
 - HTTP traffic with HTML and image content
 - FTP traffic with binary data
 - NTP traffic to synchronize the time
 - Skype traffic for audio/video chat
- Addressing on application layer: URLs/URIs
- **Socket**: interface between the application and transport layer!

Application layer – some other protocols

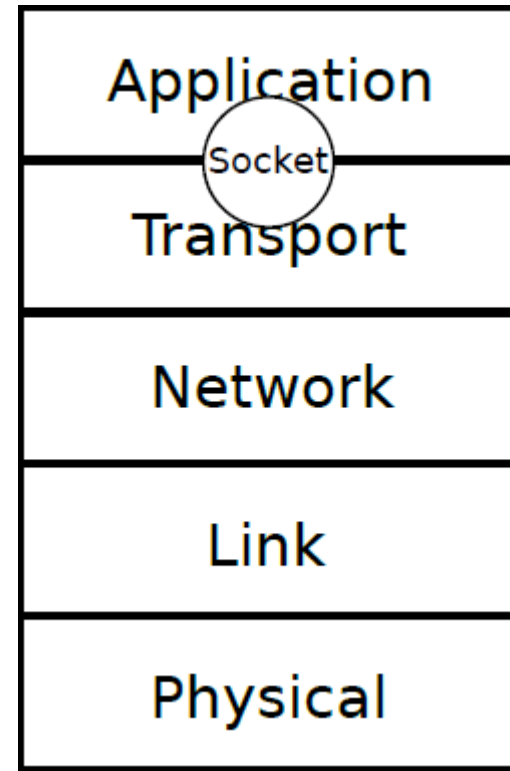
- Border Gateway Protocol (BGP)
- Domain Name Service (DNS)
- Dynamic Host Control Protocol (DHCP)
- Telnet (Teletype Network)
- File Transfer Protocol
- Internet Mail Access Protocol (IMAP)
- Simple Mail Transfer Protocol (SMTP)
- Secure Shell (SSH)
- Transport Layer Security (TLS)

Application layer addressing

- Servers can be addressed using domain names
- How can we address resources on a certain system?
- Must be globally unique name, may come with protocol identifier (allows client to use correct application to handle URL content)
- This is what uniform resource locators (URLs) provide:
 - `protocol://domain.name/resource/path`
 - `https://edimension.sutd.edu.sg/webapps/login/`
 - Globally unique, resolves to different IPs from in/outside SUTD
- Custom protocol IDs are possible, if the client can handle them
 - example: Apple's `itms://` protocolID for the iTunes store

Sockets

- From CS perspective: object that you can write, read, or listen on
- Create the socket with lower layer information (e.g., transport protocol, port number, IP address)
- Receiver will run application with similar socket
- Data pushed into sender socket will magically appear at receiver socket



Example app: TCP server

Python TCPServer

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(('', serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while 1:
    connectionSocket, addr = serverSocket.accept()

    sentence = connectionSocket.recv(1024)
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence)
    connectionSocket.close()
```

Example app: TCP client

Python TCPClient

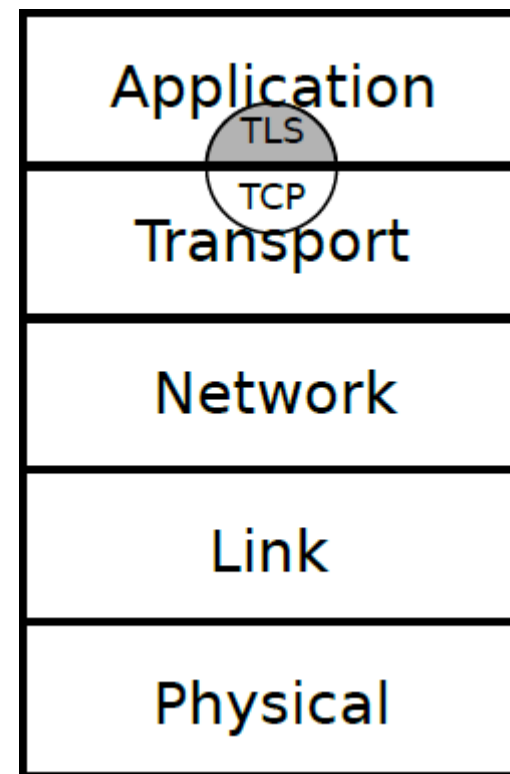
```
from socket import *
serverName = 'localhost'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence)
modifiedSentence = clientSocket.recv(1024)
print 'From Server:', modifiedSentence
clientSocket.close()
```

“Services” provided to the lower layers

- "Nice to have" properties for message transport
 - Message integrity checks
 - Good throughput
 - Low transmission delays
 - Security
- This is what is available:
 - TCP (on transport layer) provides integrity and flow control
 - UDP (on transport layer) does not provide any of the properties
- Add "security" through TLS on application layer (over TCP)
- No guarantees on low transmission delays or good throughput

Transport Layer Security (TLS)

- We will discuss TLS in detail in 50.020
- TLS is somewhat "in between" application and transport layer
- It relies on TCP, but provides services to applications
- TLS can be configured in different ways, able to provide
 - Message authentication, confidentiality, integrity
- TLS is often used to improve security of application layer protocols
 - e.g. HTTP over TLS = HTTPS




Application Layer Protocols

What are protocols?


- In general: "**language**" spoken between machines
- Protocols generally define four aspects:
 - **Types** of messages exchanged (e.g., request, response)
 - Message **syntax** (e.g. format)
 - Message **semantics** (how to interpret)
 - **Rules** for processing of messages (how to react)
- **Open** protocols have these specified publicly, but proprietary protocols also exist (e.g. Skype)
- If protocols do not encrypt the messages, it is often possible to **reverse-engineer** unspecified protocols

FTP and HTTP

FTP: separate control, data connections




FTP: separate control, data connections



- FTP client contacts FTP server at port 21, specifying TCP as transport protocol
- Client obtains authorization over control connection
- Client browses remote directory by sending commands over control connection.
- When server receives a command for a file transfer, the server opens a TCP data connection to client
- After transferring one file, server closes connection.

FTP: separate control, data connections

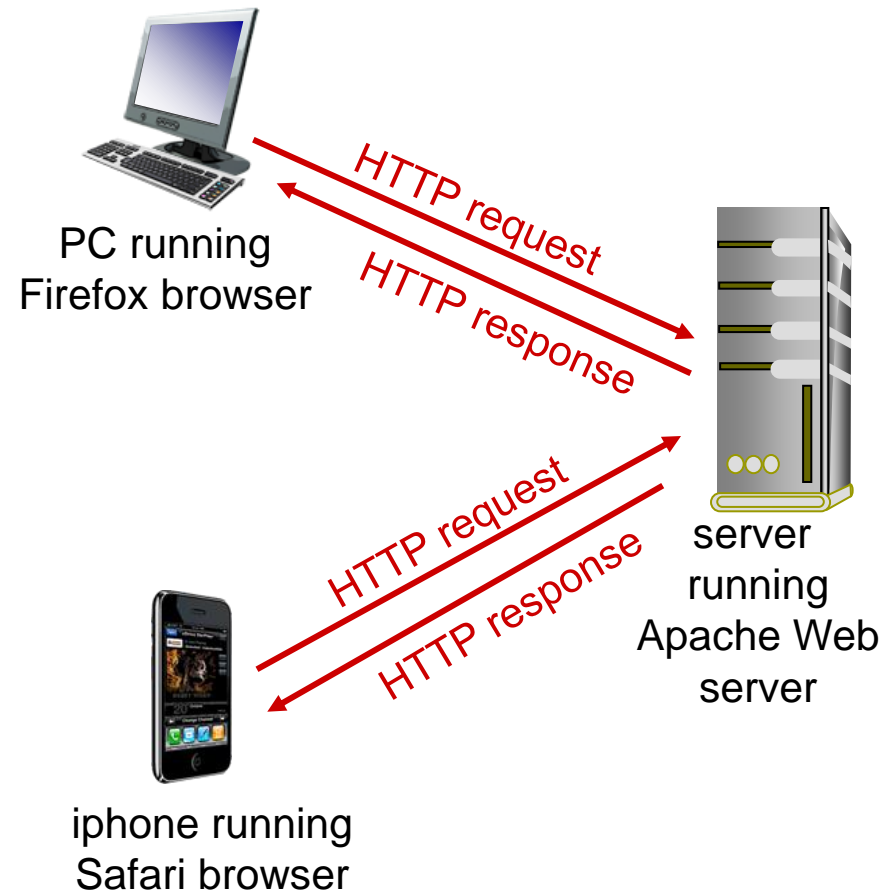


- Server opens a second TCP data connection to transfer another file.
- Control connection: “out of band”
- FTP server maintains “state”: current directory, earlier authentication

HTTP overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - *client*: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - *server*: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains no information about past client requests

aside

protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections

non-persistent HTTP

- at most one object sent over TCP connection
 - connection then closed
- downloading multiple objects required multiple connections

persistent HTTP

- multiple objects can be sent over single TCP connection between client, server

HTTP - details

- HTTP is a protocol that allows you to request files (e.g. html files, images)
 - Did you know that you can send data as well, or delete data?
- Types of messages (HTTP "verbs"): GET, PUT, POST, DELETE, HEAD, OPTIONS, CONNECT, . . .
- Message syntax: HTTP GET is ASCII-based, POST can also use binary data
 - Message start with verb, then the relative address of resources, then protocol version
- Most common message types: GET and POST
- Semantics (interpretation) and rules for processing also defined by RFC

HTTP GET

- HTTP GET is the most common verb: the client requests the server to send a resources identifies by an URL

- Example:

```
GET /demo.asp?name1=value1&name2=value2 HTTP/1.1  
host: www.sutd.edu.sg
```

- GET requests are defined to have no side-effect on the server
- Query strings can be sent with request

HTTP POST

- HTTP POST is used to submit data, e.g., forms: the client sends the server some content
- Data can be binary or string, no length limit
- Repeating a POST message will re-post the data (possibly bad)
- Example:

```
POST /demo.asp HTTP/1.1
```

```
Host: www.sutd.edu.sg
```

```
name1=value1&name2=value2
```

- GET requests are defined to have no side-effect on the server
- Query strings can be sent with request
`/demo.asp?name1=value1&name2=value2`

HTTP responses

- Status line: first line of the HTTP response, includes:
 - A numeric status code
 - Sentence with reason
- It is up to the clients to interpret the response codes

HTTP responses

- Status line: first line of the HTTP response, includes:
 - A numeric status code
 - Sentence with reason
- It is up to the clients to interpret the response codes

Common status codes

- 200 OK
- 201 Created (e.g. after successful POST)
- 301 Moved permanently
- 404 Not Found
- 500 Server error

HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

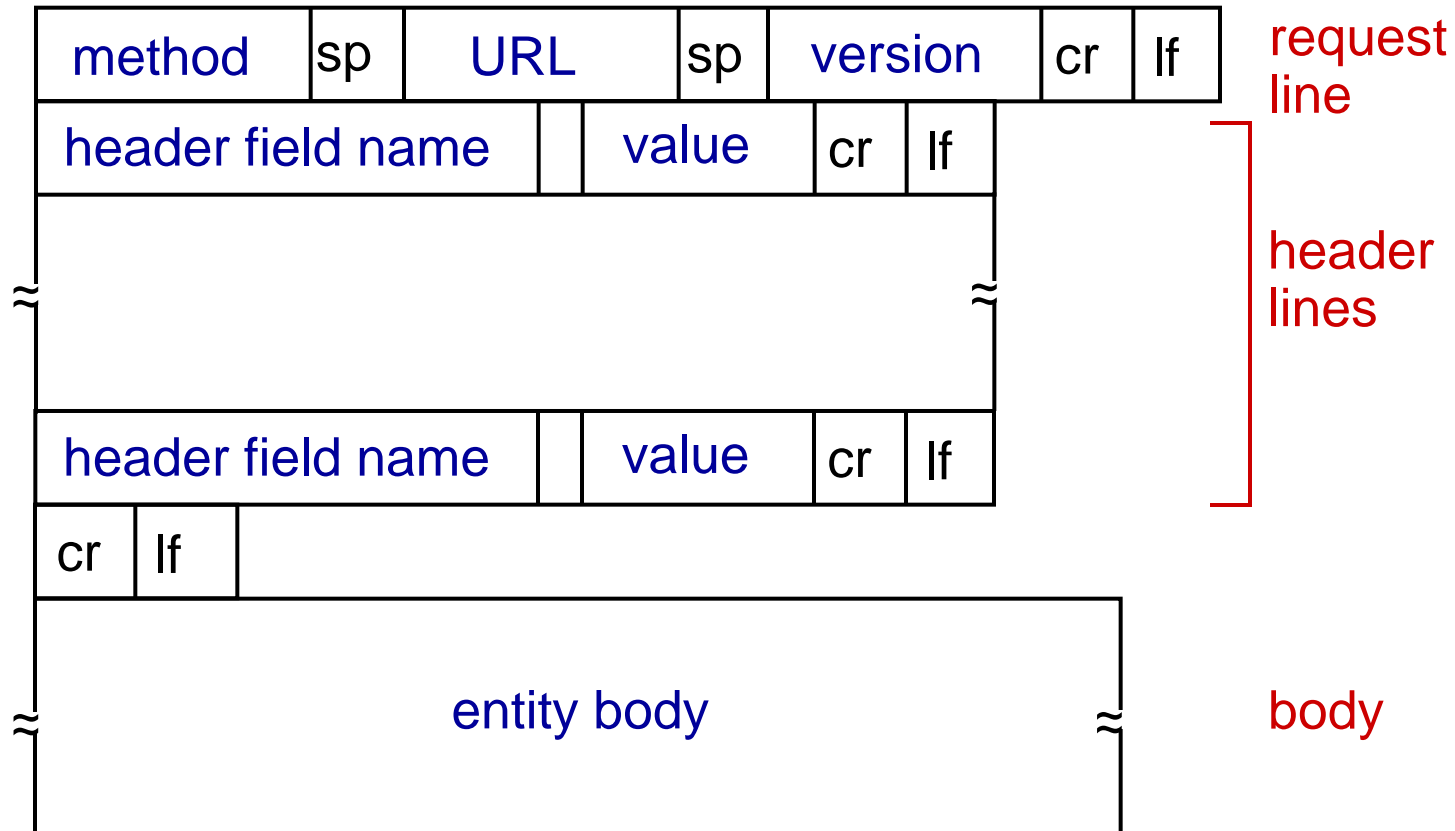
header
lines

carriage return,
line feed at start
of line indicates
end of header lines

carriage return character
line-feed character

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

HTTP request message: general format



HTTP HTML response example

```

HTTP/1.1 200 OK
Server: Apache
Content-Encoding: gzip
Content-Type: text/html; charset=utf-8
Last-Modified: Mon, 02 Sep 2016 02:20:45 GMT
Expires: Mon, 02 Sep 2016 02:21:17 GMT
Cache-Control: public, max-age=32
Content-Length: 23943
Accept-Ranges: bytes
Date: Mon, 19 Sep 2017 23:59:00 GMT
Age: 15
Connection: keep-alive
Vary: User-Agent,Accept-Encoding,X-Export-Format
  
```

status line
 (protocol
 status code
 status phrase)

header
 lines

+ all the HTML files following

Cookies

- As HTTP servers are stateless, how are sequential requests correlated?
- First connection to server creates a cookie value, subsequent communication from user includes the cookie value
- Cookies are stored in the browser




An example of P2P Applications - BitTorrent

Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)




P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks


tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Conclusion

- We discussed the application layer in more detail
- What is the interface to lower layers (Sockets)
- How addressing is done (URLs)
- Application layer protocols
- FTP and SMTP as useful application layer protocols
- HTTP as a greatly used application layer protocol
- BitTorrent as an example of P2P applications